

SARRA BEN DOUB

AI Engineer · Data Scientist · Generative AI · Document Intelligence · Machine Learning

+216 26 236 449 · sarrabendoub5@gmail.com · Tunis, Tunisia · linkedin.com/in/sarra-ben-doub · github.com/SarraBendoub · sarrabendoub.github.io/ai

PROFILE

Engineering graduate from the National School of Engineers of Tunis (ENSIT) in applied mathematics and modelling (2026), specialised in Data Science and AI, with a BSc in Data Science from the Faculty of Sciences of Tunis (FST). I build Machine Learning and LLM applications, from raw data to the tool a team uses every day: intelligent document processing (IDP), retrieval-augmented generation (RAG) and structured outputs that can be validated before anyone relies on them. Open to opportunities.

PROFESSIONAL EXPERIENCE

BIAT — Financial Document Intelligence & Credit Risk Feb – Jul 2026

Banque Internationale Arabe de Tunisie · AI Engineer & Data Scientist — Final-year project

- **Document validation & extraction.** Checking the nature of the document, then extracting financial tables from the statements filed as PDFs with the CMF using Docling and OCR (balance sheet, income statement, cash-flow statement), structured and stored in MongoDB.
- **Data cleaning & verification.** Cleaning, normalising and checking the extracted data (OCR errors, inconsistent values), then assigning it to the chart of accounts.
- **Analysis and LLM-generated summary.** Streamlit application (liquidity, solvency, profitability and leverage ratios) then a written summary with levels of assessment, the decision remaining with the analyst. Pipeline orchestrated with n8n, containerised with Docker.

Tech: Python · Docling / OCR · MongoDB · Streamlit · LLM · n8n · Docker

GAT Assurances — Claims Intelligence & NLP Jun – Aug 2025

Insurance · AI & Full-Stack Development · Internship

- **Claims-management platform:** public website, customer area and back-office, designed with a graphic designer for the visual identity and user experience.
- **Intent classification** through a French-language rule engine (missing accents, typos), backed by zero-shot classification with BART-large-MNLI when no rule decides.

Tech: Python · Flask · Transformers · BART-large-MNLI · NLP · SQLite · HTML/CSS/JS | github.com/SarraBendoub/insurance-agency-platform

VALUE — Data Engineering & Analytics Jul 2024

Data Science · Data Engineering · Internship

- **Multi-source collection** automated with Scrapy and Selenium, cleaning and merging, then **K-Means segmentation** of 314 sites and Folium mapping.

Tech: Python · Scrapy · Selenium · Pandas · scikit-learn · K-Means · Folium | github.com/SarraBendoub/HeritageScope

BNA Bank — Customer Service Chatbot Apr – May 2023

Banking · NLP & Machine Learning · Internship

- **Intent classification** on French-language queries: NLTK preprocessing, training and comparison of two architectures, FNN and LSTM, deployed with Flask.

Tech: TensorFlow · Keras · NLTK · Python · Flask | github.com/SarraBendoub/-chatbot_BNA-main

PERSONAL PROJECTS — AI & MACHINE LEARNING

CreditDoc-AI — Extraction from scanned documents & grounded Q&A 2026

Document AI · OCR · LLM · RAG

- Image preprocessing, OCR and table reconstruction on a scanned financial statement, then extraction to **JSON validated with Pydantic**: label rules first, a local LLM for the missing fields, provenance kept.
- **Grounded RAG** over the document: answers written only from the retrieved passages, page citation, explicit refusal when the information is absent.

Tech: Python · OpenCV · OCR · Ollama · Pydantic · Streamlit

News-Radar — Semantic search & RAG chatbot 2026

RAG · pgvector · Local embeddings · FastAPI

- Scheduled, idempotent RSS ingestion (~1,000 articles), embeddings computed locally and stored in **PostgreSQL / pgvector**; RAG chatbot grounded on the corpus, **FastAPI** backend.

Tech: Python · FastAPI · PostgreSQL / pgvector · fastembed · Ollama · Docker

CreditGuard — Explainable & Monitored Credit Scoring 2026

Credit risk · Probability of default

- **LightGBM** probability-of-default model, **AUC 0.757**, **SHAP** explainability per prediction and drift monitoring (PSI, Kolmogorov-Smirnov, Chi-squared). Streamlit dashboard, Docker packaging.

Tech: Python · LightGBM · SHAP · SciPy · Streamlit · Docker | github.com/SarraBendoub/CreditGuard

FraudGuard — Fraud Detection & AML/CFT Monitoring 2026

Fraud · Financial Crime · Transaction Monitoring

- **XGBoost** fraud detection on **6.4M transactions at 0.13% fraud**: **PR-AUC 0.89**, **F1 0.82**, threshold tuned on the trade-off between missed frauds and false alerts.
- Transaction-monitoring rule engine (velocity, structuring), **BLOCK / REVIEW / ALLOW** action levels, SHAP explainability on alerts, and a **data leakage found and fixed** (PR-AUC 0.9998 obtained from features posterior to the transaction).

Tech: Python · XGBoost · SHAP · pandas · Streamlit | github.com/SarraBendoub/FraudGuard

TECHNICAL SKILLS

Data Science & AI Engineering	Credit Risk · Probability of Default · Financial Statements · Fraud Detection · Model Monitoring · Explainable AI
Machine Learning	Python · Scikit-learn · XGBoost · LightGBM · SHAP · PyTorch · TensorFlow · Statistics
Data	SQL · PostgreSQL · MongoDB · Pandas · NumPy · Web scraping (Scrapy, Selenium, BeautifulSoup)
Document AI	OCR · Docling · Document Processing · LLM
Generative AI	RAG · Transformers · LLM · Ollama · Groq
Deployment	Docker · FastAPI · Streamlit · n8n · Git

EDUCATION

Engineering degree — Applied Mathematics & Modelling, National School of Engineers of Tunis	2023 – 2026
BSc — Data Science, Faculty of Sciences of Tunis	2020 – 2023
Baccalaureate — Experimental Sciences	2020

CERTIFICATION & LANGUAGES

Microsoft Azure AI Fundamentals (AI-900)
Arabic native · French fluent · English intermediate (TOEIC) · Italian A1